

3D Spatially Embedded Networks: Verifying Coordinate Backpropagation in a Feed Forward MLP

Pascal Tohouri

14 July 2025

Abstract

The purpose of this research note is to verify coordinate backpropagation for a spatially embedded feed-forward MLP. The author applies the chain rule to obtain loss-coordinate derivatives for a L -layer, N_l -width MLP embedded in 3D Euclidean Space. Mean-squared error loss is used.

1 Setup

In a prior note, the author derived the gradient descent update via backpropagation on a single spatially embedded neuron. In a dense multilayer perceptron, each layer contains multiple neurons that propagate information in parallel.

2 Example: Feed Forward MLP




Figure 1: Dense MLP with L Layers and N Neurons per Layer

The neurons weight and sum incoming edge values, applying a differentiable activation:

$$h_{n_l}^{(l)} = \sum_{n_{l-1}} w_{n_l, n_{l-1}} h_{n_{l-1}}^{(l)} \quad (1)$$

$$a_{n_l}^{(l)} = f(h_{n_l}^{(l)}). \quad (2)$$

Activation values propagate across outgoing edges to the next layer. Once spatially embedded, the n -th node in layer l has 3D coordinates $\mathbf{z}_{n_l}^{(l)} = [z_1, z_2, z_3]$.

2.1 Parameter conditioning

Edge-weights depend on Euclidean distances between connecting nodes:

$$w_{n_l, n_{l-1}}^{(l)} = g(\|\mathbf{z}_{n_l}^{(l)} - \mathbf{z}_{n_{l-1}}^{(l-1)}\|). \quad (3)$$

2.2 Training

For loss \mathcal{L} , loss-weight derivatives are calculable via the chain rule:

$$\frac{\partial \mathcal{L}}{\partial w_{n_l, n_{l-1}}} = \delta_{n_l}^{(l)} \cdot h_{n_{l-1}}^{(l-1)} \quad (4)$$

Obtain loss-coordinate derivatives by chaining weight-coordinate derivatives:

$$\frac{\partial \mathcal{L}}{\partial z_{n_l, d}} = \sum_{n_{l-1}} \frac{\partial \mathcal{L}}{\partial w_{n_l, n_{l-1}}} \cdot \frac{\partial w_{n_l, n_{l-1}}}{\partial z_{n_l, d}} + \sum_{n_{l+1}} \frac{\partial \mathcal{L}}{\partial w_{n_{l+1}, n_l}} \frac{\partial w_{n_{l+1}, n_l}}{\partial z_{n_l, d}}, \quad (5)$$

where:

$$\frac{\partial w_{n_l, n_{l-1}}}{\partial z_{n_l}} = g'(\|\mathbf{z}_{n_l} - \mathbf{z}_{n_{l-1}}\|) \frac{z_{n_l} - z_{n_{l-1}}}{\|\mathbf{z}_{n_l} - \mathbf{z}_{n_{l-1}}\|} \quad (\text{Backward}) \quad (6)$$

$$\frac{\partial w_{n_{l+1}, n_l}}{\partial z_{n_l, d}} = g'(\|\mathbf{z}_{n_{l+1}} - \mathbf{z}_{n_l}\|) \frac{z_{n_l, d} - z_{n_{l+1}, d}}{\|\mathbf{z}_{n_{l+1}} - \mathbf{z}_{n_l}\|} \quad (\text{Forward}) \quad (7)$$

In full, the total loss-coordinate gradient is:

$$\frac{\partial \mathcal{L}}{\partial z_{n_l, d}} = \sum_{n_{l-1}} \left[\delta_{n_l}^{(l)} h_{n_{l-1}}^{(l-1)} \cdot g'_{n_l, n_{l-1}} \right] \frac{z_{n_l, d} - z_{n_{l-1}, d}}{\|\mathbf{z}_{n_l} - \mathbf{z}_{n_{l-1}}\|} + \sum_{n_{l+1}} \left[\delta_{n_{l+1}}^{(l+1)} h_{n_l}^{(l)} \cdot g'_{n_{l+1}, n_l} \right] \frac{z_{n_l, d} - z_{n_{l+1}, d}}{\|\mathbf{z}_{n_{l+1}} - \mathbf{z}_{n_l}\|}. \quad (8)$$

Given the full coordinate gradient tensor for all layers l , nodes n_l , and spatial dimensions d :

$$\frac{\partial \mathcal{L}}{\partial \mathbf{Z}} \in \mathbb{R}^{L \times N_l \times D}, \quad (9)$$

apply the gradient descent update step with learning rate η :

$$z_{n_l, d}^{(l)} \leftarrow z_{n_l, d}^{(l)} - \eta \frac{\partial \mathcal{L}}{\partial z_{n_l, d}^{(l)}}. \quad (10)$$

Equivalently, using tensor notation, the update rule is succinctly expressed as:

$$\mathbf{Z} \leftarrow \mathbf{Z} - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{Z}}. \quad (11)$$

3 Conclusion

This research note presents the derivation of the gradient descent update condition for a spatially embedded feed-forward network with parameter conditioning. This method is generally applicable to feed-forward MLPs, but the note offers no empirical demonstration of the method’s efficacy. Further work will empirically evaluate the method’s efficiency, stability, and convergence.

A Notation

A.1 Function Objects

Neurons process input-output information:

$$\mathbf{X} = \begin{bmatrix} x_{1,1} & \cdots & x_{1,N_1} \\ \vdots & \ddots & \vdots \\ x_{N,1} & \cdots & x_{N,N_1} \end{bmatrix} \quad (12)$$

$$\mathbf{Y} = \begin{bmatrix} y_{1,1} & \cdots & y_{1,N_L} \\ \vdots & \ddots & \vdots \\ y_{N,1} & \cdots & y_{N,N_L} \end{bmatrix} \quad (13)$$

where:

- n is the n -th data point $n \in \{1, \dots, N\}$
- n_1 is the n -th input feature (1st layer) $n_1 \in \{1, \dots, N_1\}$
- n_L is the n -th output feature (L-th layer) $n_L \in \{1, \dots, N_L\}$

Internal functions produce hidden states and activations:

$$\mathbf{h}^{(l)} = \begin{bmatrix} h_1 \\ \vdots \\ h_{N_l} \end{bmatrix}, \quad \mathbf{a}^{(l)} = \begin{bmatrix} a_1 \\ \vdots \\ a_{N_l} \end{bmatrix} \quad (14)$$

where:

- n_l is the n -th internal node in the l -th layer $n_l \in \{1, \dots, N_l\}$

A.2 Spatial Objects

Distinguish function objects from spatial coordinates. Here is a ‘slice’ of the network at the l -th layer:

$$\mathbf{Z}^{(l)} = \begin{bmatrix} z_{1,1}^{(l)} & \cdots & z_{1,D}^{(l)} \\ \vdots & \ddots & \vdots \\ z_{N_l,1}^{(l)} & \cdots & z_{N_l,D}^{(l)} \end{bmatrix}_{N_l \times D} \quad (15)$$

where:

- $z_{n_l,d}^{(l)}$ is the n_l -th node’s d -th dimension coordinate $d \in \{1, \dots, D\}$, for the l -th layer $l \in \{1, \dots, L\}$